

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, roc_auc_score, roc_curve

from sklearn.cluster import KMeans
from mlxtend.frequent_patterns import apriori, association_rules

import warnings
warnings.filterwarnings('ignore', category=DeprecationWarning)

```

```

churn_df = pd.read_csv('/content/WA_Fn-UseC_-Telco-Customer-Churn.csv')
retail_df = pd.read_excel('/content/online_retail_II.xlsx')

```

```

print("Churn Data Sample:")
display(churn_df.head())

```

```

print("\nOnline Retail Sample:")
display(retail_df.head())

```

Churn Data Sample:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...
0	7590-VHVEG	Female		0	Yes	No	1	No	No phone service	DSL	No ..
1	5575-GNVDE	Male		0	No	No	34	Yes	No	DSL	Yes ..
2	3668-QPYBK	Male		0	No	No	2	Yes	No	DSL	Yes ..
3	7795-CFOCW	Male		0	No	No	45	No	No phone service	DSL	Yes ..
4	9237-HQITU	Female		0	No	No	2	Yes	No	Fiber optic	No ..

5 rows × 21 columns

Online Retail Sample:

Invoice	StockCode	Description	Quantity	InvoiceDate	Price	Customer ID	Country
0	489434	85048 15CM CHRISTMAS GLASS BALL 20 LIGHTS	12	2009-12-01 07:45:00	6.95	13085.0	United Kingdom
1	489434	79323P PINK CHERRY LIGHTS	12	2009-12-01 07:45:00	6.75	13085.0	United Kingdom
2	489434	79323W WHITE CHERRY LIGHTS	12	2009-12-01 07:45:00	6.75	13085.0	United Kingdom
3	489434	22041 RECORD FRAME 7" SINGLE SIZE	48	2009-12-01 07:45:00	2.10	13085.0	United Kingdom
4	489434	21232 STRAWBERRY CERAMIC TRINKET BOX	24	2009-12-01 07:45:00	1.25	13085.0	United Kingdom

▼ Question 1 – Customer Churn Prediction

Data Preprocessing

```

churn_df['TotalCharges'] = churn_df['TotalCharges'].replace(' ', np.nan)
churn_df.dropna(subset=['TotalCharges'], inplace=True)
churn_df['TotalCharges'] = churn_df['TotalCharges'].astype(float)

```

```
for col in churn_df.columns:
    if churn_df[col].dtype == 'object':
        le = LabelEncoder()
        churn_df[col] = le.fit_transform(churn_df[col])

X = churn_df.drop('Churn', axis=1)
y = churn_df['Churn']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Model Training (Random Forest)

```
rf_model = RandomForestClassifier(n_estimators=200, random_state=42)
rf_model.fit(X_train, y_train)
y_pred = rf_model.predict(X_test)
```

Evaluation + Visuals

```
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))


cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix - Customer Churn')
plt.xlabel('Predicted'); plt.ylabel('Actual'); plt.show()

y_pred_proba = rf_model.predict_proba(X_test)[:,1]
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="AUC="+str(round(roc_auc_score(y_test, y_pred_proba),2)))
plt.legend(); plt.title('ROC Curve - Churn Prediction'); plt.show()
```

```
Accuracy: 0.7917555081734187
```


Classification Report:				
	precision	recall	f1-score	support
0	0.83	0.91	0.86	1033
1	0.65	0.48	0.55	374
accuracy			0.79	1407
macro avg	0.74	0.69	0.71	1407
weighted avg	0.78	0.79	0.78	1407

Confusion Matrix - Customer Churn



```
importances = pd.Series(rf_model.feature_importances_, index=X.columns).sort_values(ascending=False)[:10]
importances.plot(kind='barh')
plt.title('Top 10 Feature Importances - Churn Model')
plt.show()
```

Top 10 Feature Importances - Churn Model



▼ Question 2 – Customer Segmentation (Online Retail)

Data Cleaning

```
print(retail_df.info())
print(retail_df.head())

retail_df = retail_df.dropna(subset=['Customer ID', 'Description'])

retail_df['Customer ID'] = retail_df['Customer ID'].astype(int)

retail_df = retail_df[retail_df['Quantity'] > 0]
```

```
totalprice = retail_df['Quantity'] * retail_df['Price']
```

```
print("After cleaning:")
print(retail_df.describe(include='all'))
print(retail_df.isnull().sum())
```

2	489434	79323W	WHITE CHERRY LIGHTS	12	
3	489434	22041	RECORD FRAME 7" SINGLE SIZE	48	
4	489434	21232	STRAWBERRY CERAMIC TRINKET BOX	24	

0	2009-12-01 07:45:00	6.95	13085	United Kingdom	83.4
1	2009-12-01 07:45:00	6.75	13085	United Kingdom	81.0
2	2009-12-01 07:45:00	6.75	13085	United Kingdom	81.0
3	2009-12-01 07:45:00	2.10	13085	United Kingdom	100.8
4	2009-12-01 07:45:00	1.25	13085	United Kingdom	30.0

After cleaning:


count	407695.0	407695		Description	Quantity \
unique	19215.0	4017		407695	407695.000000
top	500356.0	85123A	WHITE HANGING HEART T-LIGHT HOLDER		NaN
freq	270.0	3153		3153	NaN
mean	Nan	Nan		Nan	13.586686
min	Nan	Nan		Nan	1.000000
25%	Nan	Nan		Nan	2.000000
50%	Nan	Nan		Nan	5.000000
75%	Nan	Nan		Nan	12.000000
max	Nan	Nan		Nan	19152.000000
std	Nan	Nan		Nan	96.842229

count		407695	407695.000000	407695.000000	\
unique		Nan	Nan	Nan	
top		Nan	Nan	Nan	
freq		Nan	Nan	Nan	
mean	2010-07-01 10:10:10.782177792		3.294188	15368.504107	
min	2009-12-01 07:45:00		0.000000	12346.000000	
25%	2010-03-26 14:01:00		1.250000	13997.000000	
50%	2010-07-08 15:46:00		1.950000	15321.000000	
75%	2010-10-14 17:09:00		3.750000	16812.000000	
max	2010-12-09 20:01:00		10953.500000	18287.000000	
std		Nan	34.756655	1679.795700	

count		407695	407695.000000		
unique		37	Nan		
top	United Kingdom		Nan		
freq	370951		Nan		
mean	Nan	21.663261			
min	Nan	0.000000			
25%	Nan	4.950000			
50%	Nan	11.900000			
75%	Nan	19.500000			
max	Nan	15818.400000			
std	Nan	77.147356			
Invoice	0				
StockCode	0				
Description	0				
Quantity	0				
InvoiceDate	0				
Price	0				
Customer ID	0				
Country	0				
totalprice	0				

```
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10,5))
sns.histplot(retail_df['totalprice'], bins=50, kde=True)
plt.title('Distribution of Total Transaction Values')
plt.xlabel('Total Price per Invoice Line')
plt.ylabel('Count')
plt.show()
```



```
rfm = retail_df.groupby('customerid').agg({
    'invoicedate': lambda x: (retail_df['invoicedate'].max() - x.max()).days,
    'invoice': 'count',
    'totalprice': 'sum'
}).rename(columns={'invoicedate': 'recency', 'invoice': 'frequency', 'totalprice': 'monetary'})

rfm.reset_index(inplace=True)
rfm.head()
```

	Customer ID	recency	frequency	monetary	
0	12346	164	33	372.86	
1	12347	2	71	1323.32	
2	12348	73	20	222.16	
3	12349	42	102	2671.14	
4	12351	10	21	300.93	

Next steps: [Generate code with rfm](#) [New interactive sheet](#)

: Scale & Cluster (K-Means)

```
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans

scaler = StandardScaler()
rfm_scaled = scaler.fit_transform(rfm[['recency', 'frequency', 'monetary']])

kmeans = KMeans(n_clusters=4, random_state=42)
rfm['cluster'] = kmeans.fit_predict(rfm_scaled)

print(rfm.groupby('cluster').mean())
```

cluster	Customer ID	recency	frequency	monetary
0	15334.840728	43.057772	84.812959	1592.729466
1	15403.817308	243.004808	29.875962	614.363956
2	15249.846715	13.000000	732.014599	16472.760460
3	15453.750000	3.750000	2654.750000	236568.790000


: Visualize Clusters

```

import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(8,6))
sns.scatterplot(data=rfm, x='recency', y='monetary', hue='cluster', palette='Set2', s=80)
plt.title('Customer Segmentation using K-Means (RFM Model)')
plt.xlabel('Recency (Days)')
plt.ylabel('Monetary (Total Spend)')
plt.legend(title='Cluster')
plt.show()

```



❖ Question 3 – Cross-Selling Opportunities

Prepare Basket Data

◆ Gemini

```

basket = (retail_df
    .groupby(['Invoice', 'description'])['quantity']
    .groupby(['Invoice', 'Description'])['Quantity']
    .sum().unstack().fillna(0))

basket = basket.applymap(lambda x: 1 if x > 0 else 0)

print("Basket data shape:", basket.shape)
basket.head()

```

```
/tmp/ipython-input-2011280971.py:6: FutureWarning: DataFrame.applymap has been deprecated. Use DataFrame.map instead.
  basket = basket.applymap(lambda x: 1 if x > 0 else 0)
Basket data shape: (19215, 4444)
```

Description	DOORMAT UNION JACK GUNS AND ROSES	3 STRIPEY MICE FELTCRAFT	PURPLE FLOCK DINNER CANDLES	4 ANIMAL STICKERS	BLACK PIRATE TREASURE CHEST	BROWN PIRATE TREASURE CHEST	Bank Charges	CAMPHOR WOOD PORTOBELLO MUSHROOM	CHERRY BLOSSOM DECORATIVE FLASK	FAIRY CAKE CANDLES	... LATTICE CHARGER LARGE	ZINC HEART HE CHAR SM
-------------	--	--------------------------------	--------------------------------------	-------------------------	--------------------------------------	--------------------------------------	-----------------	---	--	--------------------------	------------------------------------	-----------------------------------

Association Rule Mining

```
frequent_items = apriori(basket, min_support=0.02, use_colnames=True)
rules = association_rules(frequent_items, metric="lift", min_threshold=1)
rules.sort_values('lift', ascending=False).head(10)
```

	409437	0	0	0	0	0	0	0	0	0	0	0
	489420	precedents	0	consequents	antecedent support	consequent support	Support	Confidence	lift	representativity	leverage	conviction
28	(WOODEN FRAME ANTIQUE WHITE)	(WOODEN PICTURE FRAME WHITE FINISH)	0.052511	0.042571	0.028832	0.549058	12.897504			1.0	0.026596	2.123178
29	(WOODEN PICTURE FRAME WHITE FINISH)	(WOODEN PICTURE ANTIQUE WHITE)	0.042571	0.052511	0.028832	0.677262	12.897504			1.0	0.026596	2.935780
25	(SWEETHEART CERAMIC TRINKET BOX)	(STRAWBERRY CERAMIC TRINKET BOX)	0.042050	0.069477	0.032371	0.769802	11.079959			1.0	0.029449	4.042272
24	(STRAWBERRY CERAMIC TRINKET BOX)	(SWEETHEART CERAMIC TRINKET BOX)	0.069477	0.042050	0.032371	0.465918	11.079959			1.0	0.029449	1.793636
7	(CHOCOLATE HOT WATER BOTTLE)	(HOT WATER BOTTLE TEA AND SYMPATHY)	0.041686	0.043768	0.020088	0.481898	11.010301			1.0	0.018264	1.845643
6	(HOT WATER BOTTLE TEA AND SYMPATHY)	(CHOCOLATE HOT WATER BOTTLE)	0.043768	0.041686	0.020088	0.458977	11.010301			1.0	0.018264	1.771301
	(HEART OF	(HEART OF										

Visualization

```
plt.figure(figsize=(8,6))
sns.scatterplot(x='support', y='confidence', size='lift', data=rules)
plt.title('Association Rules - Cross-Selling Opportunities')
plt.show()
```



Combined Summary Output


```
summary = pd.DataFrame({
    'Question': ['Churn Prediction', 'Customer Segmentation', 'Cross-Selling'],
    'Algorithm': ['Random Forest', 'K-Means', 'Apriori'],
    'Goal': ['Predict churn likelihood', 'Group customers by RFM behavior', 'Find product associations'],
    'Key Metric': ['Accuracy/AUC', 'Cluster Means', 'Lift/Support']
})
display(summary)
```

	Question	Algorithm	Goal	Key Metric	Actions
0	Churn Prediction	Random Forest	Predict churn likelihood	Accuracy/AUC	
1	Customer Segmentation	K-Means	Group customers by RFM behavior	Cluster Means	
2	Cross-Selling	Apriori	Find product associations	Lift/Support	


```
import matplotlib.pyplot as plt

churn_raw = pd.read_csv('/content/WA_Fn-UseC_-Telco-Customer-Churn.csv')

plt.figure(figsize=(7,5))
sns.barplot(data=churn_raw, x='Contract', y=churn_raw['Churn'].apply(lambda x: 1 if x=='Yes' else 0))
plt.title('Churn Rate by Contract Type')
plt.ylabel('Churn Rate')
plt.xlabel('Contract Type')
plt.show()
```



```
plt.figure(figsize=(6,4))
sns.heatmap(rfm[['recency','frequency','monetary']].corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation between RFM Metrics')
plt.show()
```



Top Product Co-Occurrences Network (Cross-Selling)

```
import networkx as nx

top_rules = rules.sort_values('lift', ascending=False).head(10)

G = nx.Graph()
for _, row in top_rules.iterrows():
    for a in row['antecedents']:
        for b in row['consequents']:
            G.add_edge(a, b, weight=row['lift'])

plt.figure(figsize=(10,7))
pos = nx.spring_layout(G, k=0.5)
nx.draw(G, pos,
        with_labels=True,
        node_size=2000,
        node_color='lightgreen',
        font_size=10,
        font_weight='bold',
        edge_color='gray')
plt.title('Top Product Associations Network (Cross-Selling)')
plt.show()
```

Top Product Associations Network (Cross-Selling)

